Junyoung Park

AI Engineer | GPU Inference and Multimodal Training

junpark314a@gmail.com | +82-10-7709-2772 | GitHub | LinkedIn

Education

Mar 2018 - Aug 2024

Seoul National University

B.S. in Computer Science and Engineering | GPA: 3.51 / 4.3

Experience

Apr 2026 - Present

Nari Labs | AI Engineer

- Developed and trained a causal avatar video model conditioned on text, images and audio.
- Built optimized inference engines for Qwen3-TTS and full-duplex speech models; developed and operated a streaming TTS API.

Aug 2023 - Jun 2025

Weavel | Co-founder & CTO

- Built Promptmodel (prompt versioning), Weavel Analytics (LLM logs and evaluation), Ape (automated prompt optimization), and Typa (an AI agent for LinkedIn marketing).
- Raised pre-seed funding from Krew Capital; selected for Y Combinator S24.

Selected Projects

Aug - Sep 2026

Full-duplex Speech-to-Speech Inference Engines

- Rebuilt the runtime and scheduler for full-duplex models around 80 ms audio deadlines, overcoming inefficient batching in vLLM.
- Served 80 concurrent PersonaPlex 7B sessions or 64 Nemotron VoiceChat 11B sessions on one H100; expanded Nemotron capacity to 16x NVIDIA NIM.

Aug - Sep 2026

Qwen3-TTS Inference Engine and Streaming API

- Coordinated work across Qwen3-TTS submodules in a custom scheduler; built and operated a streaming TTS API.
- Ranked 1st on Speko and 2nd on Coval for time to first audio (Aug - Sep 2026).

Apr - Jul 2026

2B Causal Avatar Model and Training Optimization

- Built a text-, image- and audio-conditioned video model through teacher training, ODE distillation, self-forcing DMD, and adversarial refinement.
- Optimized FSDP2 training: reduced teacher step time by 37% and increased model FLOPs utilization from 23.3% to 37.0%.

Oct - Nov 2025

Async-GRPO Weight Synchronization Optimization

- Extended Red Hat's open-source GRPO framework: replaced Ray Object Store transfers between trainer and rollout workers with direct GPU communication via NCCL.
- Reduced 7B model weight synchronization from 224.06 s to 0.256 s (874x faster).

Research

2024

Findings of EMNLP 2024 | Co-author

Towards Efficient Visual-Language Alignment of the Q-Former for Visual Reasoning Tasks

- Analyzed parameter-efficient tuning and sublayer importance in InstructBLIP for visual reasoning.

Awards & Activities

2024

Y Combinator S24

Selected as a Weavel co-founder.

Sep 2023

OpenAI x SKT Prompter Day Seoul | Excellence Award

Built a multilingual interactive storytelling product for children.

Technical Skills Python, PyTorch, FSDP2, NCCL, Ray, CUDA Graphs; inference scheduling, streaming APIs, multimodal training and distillation